

Vector:

Vectors are straightforward to implement and offer fast access times due to their contiguous memory storage. The time used for reading and going through data is $O(n)$, and inserting elements is typically $O(1)$ on average. However, vectors can become inefficient due to the potential overhead of resizing, leading to a worst-case complexity of $O(n)$ for insertions. Vectors are not well-suited for frequent insertions or deletions, which could lead to performance issues as the dataset grows.

Vectors	Line Cost	# Times Executes	Total Cost
Opening and reading a file	1	$O(n)$	$O(n)$
Each line and Creating courses Objects	1	$O(n)$	$O(n)$

Hash Table:

Hash tables provide excellent average-case performance. Reading and parsing data remains $O(n)$, while insertion into the hash table is $O(1)$ on average. The worst-case complexity for hash tables can degrade to $O(n^2)$ if many collisions occur, although this is highly unlikely with a good hash function. Hash tables are memory efficient and offer rapid lookup, insertion, and deletion times, making them a strong candidate for this project. The primary drawback lies in managing collisions, which can complicate the implementation.

Hash Table	Line Cost	# Times Executes	Total Cost
Opening and reading a file	1	$O(n)$	$O(n)$
Each line and Creating courses Objects	$O(1)$	$O(n)$	$O(n)$

Binary Search Tree:

Binary search trees offer efficient search, insert, and delete operations, especially when balanced, with an average-case time complexity of $O(\log n)$. The worst-case problem can degrade to $O(n^2)$ if the tree becomes unbalanced, leading to inefficiencies. Trees also require additional memory to store pointers and come with the added complexity of maintaining balance.

Binary Search Tree	Line Cost	# Times Executes	Total Cost
Opening and reading a file	1	$O(n)$	$O(n)$
Each line and Creating courses Objects	$O(\log n)$	$O(n)$	$O(n \log n)$

Based on the analysis, I recommend using a hash table for this project. Hash tables offer the best average-case performance for operations such as insertions, lookups, and deletions, all of which run in $O(1)$ time. They are also efficient in terms of memory usage, particularly when appropriately sized. While hash tables have a worst-case performance of $O(n)$ due to potential collisions, this risk can be minimized with a well-designed hash function and careful load factor management.

In contrast, although a tree structure could provide more consistent performance if balanced, the complexity and overhead involved in maintaining that balance may not be justified for this specific application. A vector is straightforward to implement but may run into inefficiencies due to resizing and slower lookups as the dataset grows.

